

PROGRAM

SYNTHESIS

Nathanaël Fijalkow

CNRS, LaBRI, Bordeaux, France

& University of Warsaw, Poland

OUTLINE

- The multiple benefits of Domain Specific Language

"a program has structure, it is not a raw piece of text"

- The need for proper interfaces

"LLMs + fact-checking = future"

- The **TRUE** Programmatic Reinforcement Learning

"Program you are, and to program you will return"

environment

policy

BACK IN 2011: FLASHFILL (Microsoft)

The screenshot shows a Microsoft Excel spreadsheet titled "Summit Attendees". The data consists of two columns: "Attendees" and "Office Address". In column C, there is a formula bar with the text "Hi _____" and a dropdown arrow. The formula has been applied to the cell, resulting in the text "Hi Gerald" in the adjacent cell. The rest of the table contains names and addresses from various locations like Chicago, New York, and Boulder.

A	B	C	D	E	F
1 Attendees	Office Address	Hi _____	E-mail	Zip Code	
2 Gerald Parker	320 N Morgan St #600, Chicago, IL 60607	Hi Gerald			
3 Phillip Davis	111 8th Ave. New York, NY 10011				
4 Ralph Gonzales	2930 Pearl Street, Boulder, CO. 80301				
5 Shawn Jenkins	1160 Bordeaux Drive, Sunnyvale, CA 94089				
6 Thomas Murphy	1160 Bordeaux Drive, Sunnyvale, CA 94089				
7 Norma Long	2930 Pearl Street, Boulder, CO. 80301				
8 Beverly Harris	320 N Morgan St #600, Chicago, IL 60607				
9 Melissa Torres	1160 Bordeaux Drive, Sunnyvale, CA 94089				
10 Joshua Watson	111 8th Ave. New York, NY 10011				
11 Larry Wright	500 W 2nd St. Austin, TX 78701				
12 Lois Brooks	2930 Pearl Street, Boulder, CO. 80301				
13 Steve James	1160 Bordeaux Drive, Sunnyvale, CA 94089				
14 Helen Rivera	2930 Pearl Street, Boulder, CO. 80301				
15 Samuel Bailey	111 8th Ave. New York, NY 10011				
16 Kenneth Gonzalez	2930 Pearl Street, Boulder, CO. 80301				
17 Sharon Washington	1160 Bordeaux Drive, Sunnyvale, CA 94089				
18 Mildred Evans	1160 Bordeaux Drive, Sunnyvale, CA 94089				
19					
20					
21					
22					

(this picture is actually Smartfill, Google's version)

- Limitations:
- it doesn't know that Paris is in France!
 - it doesn't output the program so we cannot trust the results

BACK IN 2021: CODEX (OpenAI / Github)

```
1 import datetime
2
3 def parse_expenses(expenses_string):
4     """Parse the list of expenses and return the list of triples (date, value, currency).
5     Ignore lines starting with #.
6     Parse the date using datetime.
7     Example expenses_string:
8         2016-01-02 -34.01 USD
9         2016-01-03 2.59 DKK
10        2016-01-03 -2.72 EUR
11    """
12    expenses = []
13    for line in expenses_string.splitlines():
14        if line.startswith("#"):
15            continue
16        date, value, currency = line.split(" ")
17        expenses.append((datetime.datetime.strptime(date, "%Y-%m-%d"),
18                         float(value),
19                         currency))
20    return expenses
```

Copilot

Limitations: treats code as raw text: a lot of syntactic errors!

A FEW MONTHS AGO: BARD (Google)

What new discoveries from the James Space Webb Telescope can I tell my 9 year old about?

Your 9-year old might like these recent discoveries made by The James Webb Space Telescope (JWST):

- In 2023, The JWST discovered a number of galaxies nicknamed "green peas." They were given this name because they are small and round, like peas.
- The telescope captured images of galaxies that are over 13 billion years old. This means that the light from these galaxies has been traveling for over 13 billion years to reach us.
- JWST took the very first pictures of a planet outside of our own solar system. These distant worlds are called "exoplanets". Exo means "from outside".

These discoveries can spark a child's imagination about the infinite wonders of the universe

Upvote Downvote Check it

⋮

Limitations: it doesn't know, it learns (from biased data)

SPECIFICATION

Program
Synthesis



PROGRAM

logical formula
examples
natural language description
demonstration

:

code
System
chip
macro

:

- One of the oldest dream of AI
- For decades only for delusional researchers
- The upcoming revolution for / developpers
end users ?

RESEARCH PERSPECTIVE ON SYNTHESIS

- **HUGE** practical success and impact
- Theory is lagging behind but **highly needed**
- at the **intersection** of
 - * Programming languages
 - * Software Engineering
 - * Machine learning
 - * Algorithms

PART I SYNTHESIS IN A DSL

All Large Language Models (LLM) tools I know write code in Python / C / Java / ...

I believe this is **WRONG**: complicated syntax / semantics

program aliasing

no partial executions

hard to reason with

no lifting to other languages

basis errors

DSL : Domain Specific Language

- Designed for a specific class of tasks
- Defined by a family of primitives

⇒ FUNCTIONAL PROGRAMMING

EXAMPLE DSL

list of primitives

sort :	$list(int) \rightarrow list(int)$
map :	$(t_0 \rightarrow t_1) \rightarrow list(t_0) \rightarrow list(t_1)$
succ :	$int \rightarrow int$
:	

their types

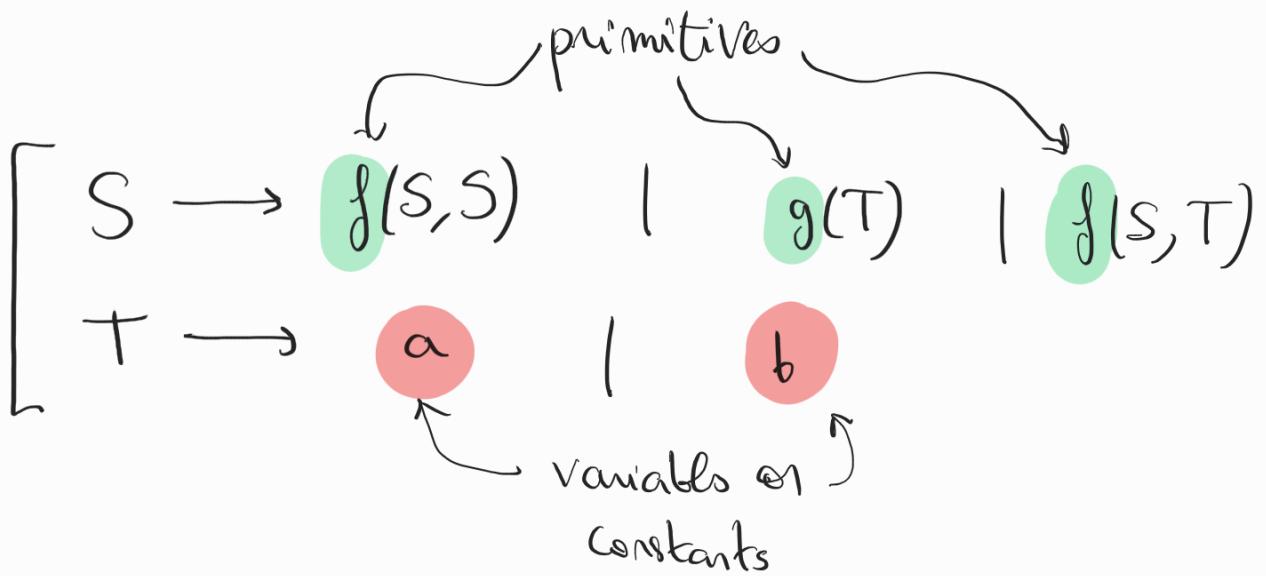
TYPES

atomic types : int, bool, string, list(int), ...

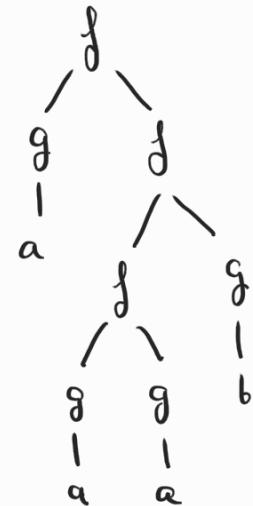
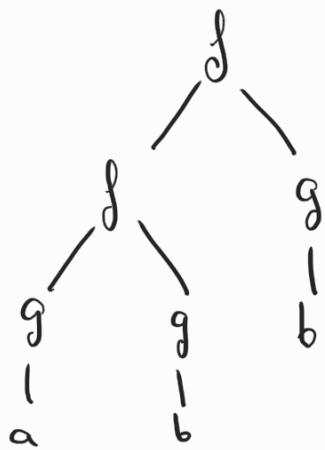
generic types : t_0, t_1, \dots

constructor : $Z_1 \rightarrow Z_2 : functions$

CONTEXT FREE GRAMMARS



Example trees generated by CFG :



$L(G)$: set of trees generated by the CFG G

SOME EXAMPLES OF REASONING



List of primitives

Add:	int → int → int
Double:	int → int
Halve:	int → int
IfThenElse:	bool → int → int → int
Even:	int → bool
Equal:	int → int → bool
0: int	False: bool
1: int	True: bool

type request:
int → int

Compilation

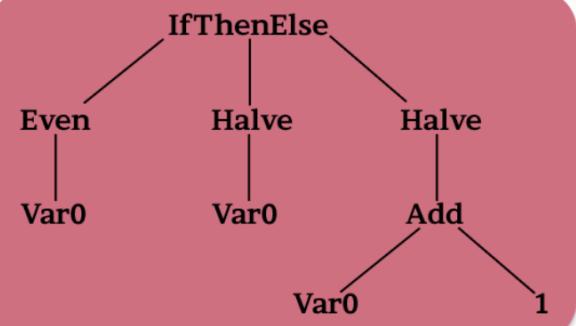
Grammar

int → Add(int, int)
int → Double(int)
int → Halve(int)
int → IfThenElse(bool, int, int)
int → 0
int → 1
int → Var0
bool → Even(int)
bool → Equal(int, int)
bool → True
bool → False

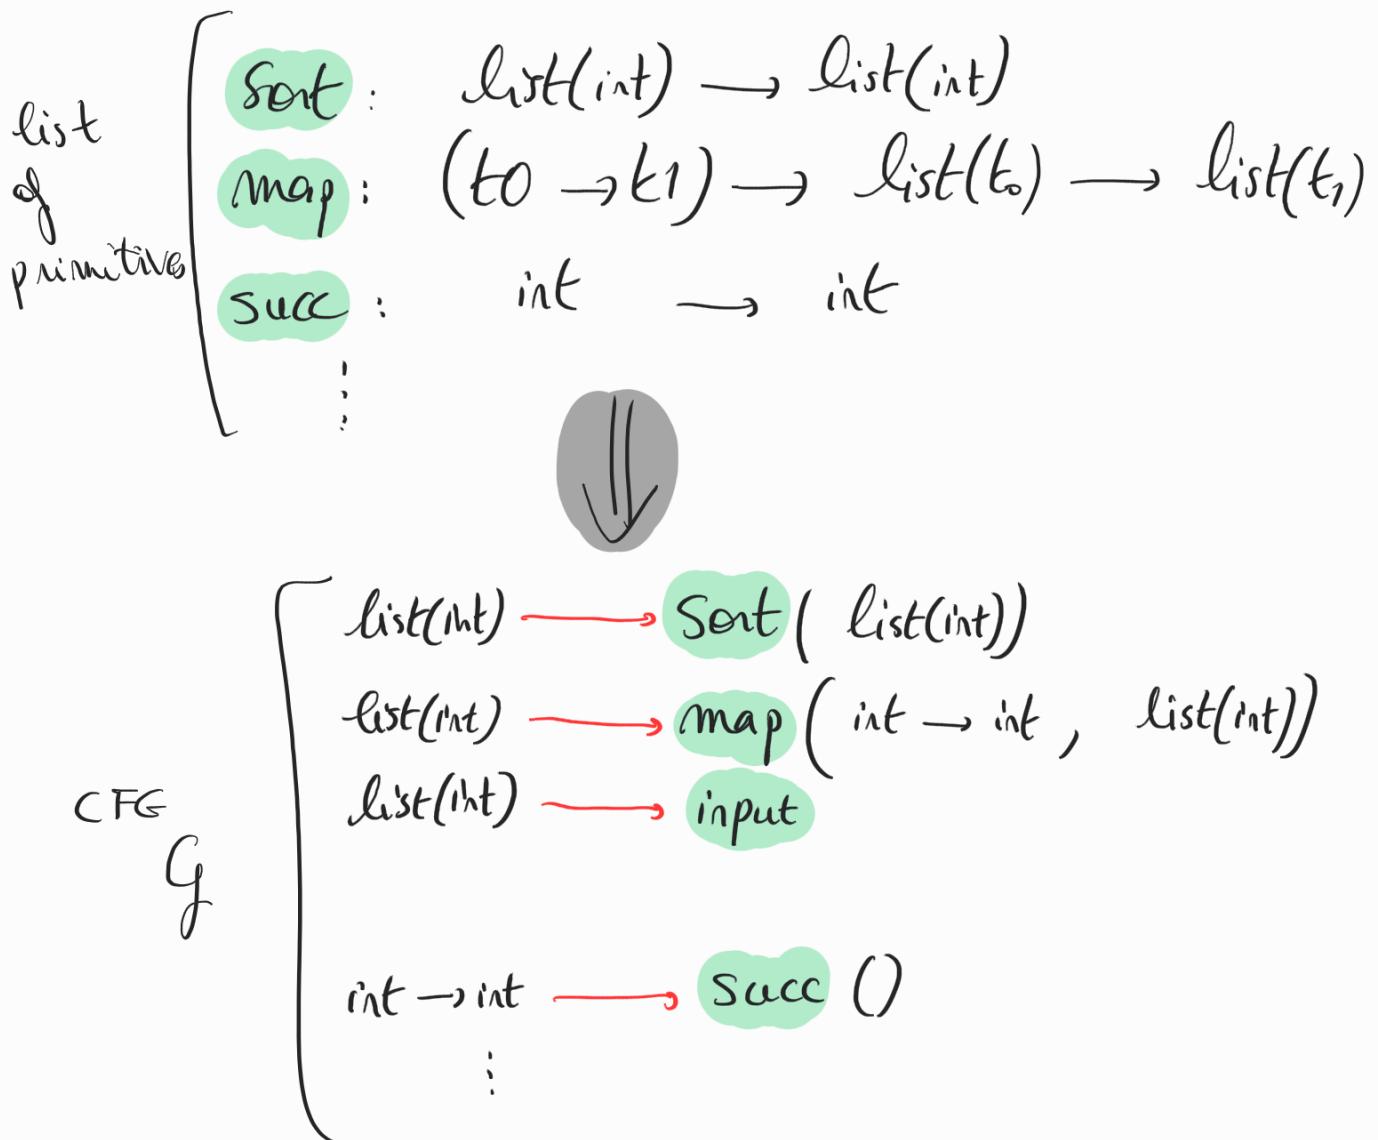
Usual (Python-style) syntax

```
if Even(var0):  
    Halve(var0)  
else:  
    Halve(Add var0 1)
```

Equivalent AST representation

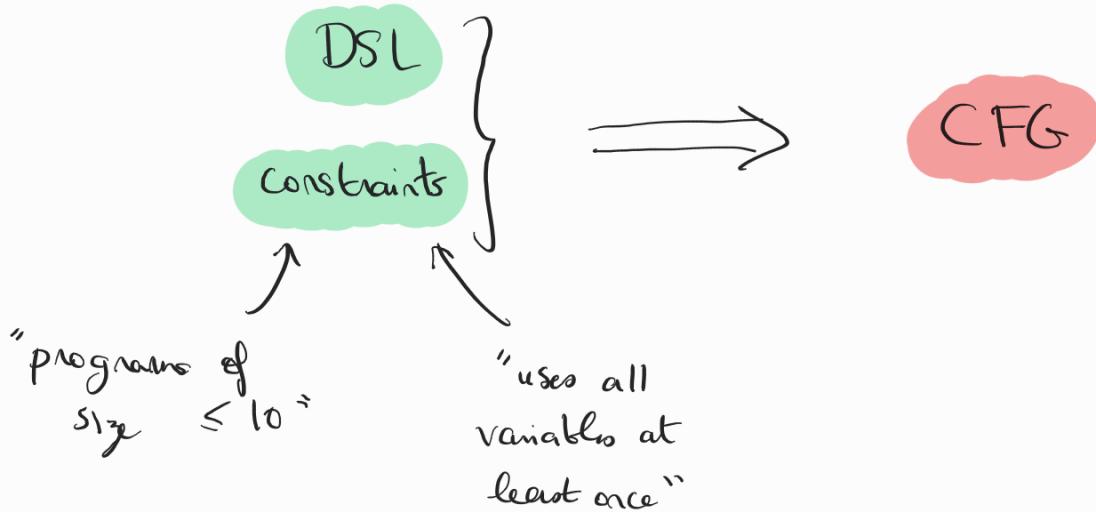


TYPE-BASED COMPIRATION



$\mathcal{L}(G) = \text{all correctly typed programs}$

ADVANCED COMPILATION



RECENT WORK: SEMANTICAL PRUNING

automatically generating equations
to reduce the DSL

(Theo Matignon)

Example equations:

- $\text{Halve}(\text{Double}(P)) \equiv P$
- $\text{Add}(P, \text{Add}(Q, R)) \equiv \text{Add}(\text{Add}(P, Q), R)$
- $\text{Add}(P, Q) \equiv \text{Add}(Q, P)$

Objective: construct a CFG representing programs
up to equivalence

Unfortunately that's impossible: equivalence of programs
is very quickly undecidable

Remark: $\text{Add}(P, Q) \equiv \text{Add}(Q, P)$ is already too hard

But: we can quotient by "regular equivalence"

Key ingredient: well order on programs

PART II KNOWLEDGE MEETS SYNTHESIS

Smartfill doesn't know that Paris is in France

Band doesn't know what James Web Telescope does

ChatGPT cannot add numbers

BUT

Wikipedia and/or Calculators do !

So far, synthesis was either:

- purely syntactic (FlashFill)
- or looks for semantics in training data

RESEARCH QUESTION

How to use external knowledge in Synthesis ?

RECENT WORK: WIKI CODER (Theo Matignon)

a program synthesis tool which queries Wikipedia when it's lost

BUT: Wikipedia's structure is hard to use

A LOT to be done

PART III

PROGRAMMATIC RL

Experiment on MountainCar

In ~20s we found : Input variables (position, velocity)

```
If ( sign(var1) ≥ var0 ) :  
    left  
Else:  
    Right
```

Input variables (position, velocity)

a program

Baseline Deep learning (DQN) : 5 mn for a neural policy ...

QUESTIONS

Q1 What is a good DSL for policies ?

Q2 how do we synthesise policies as programs ?

So far :

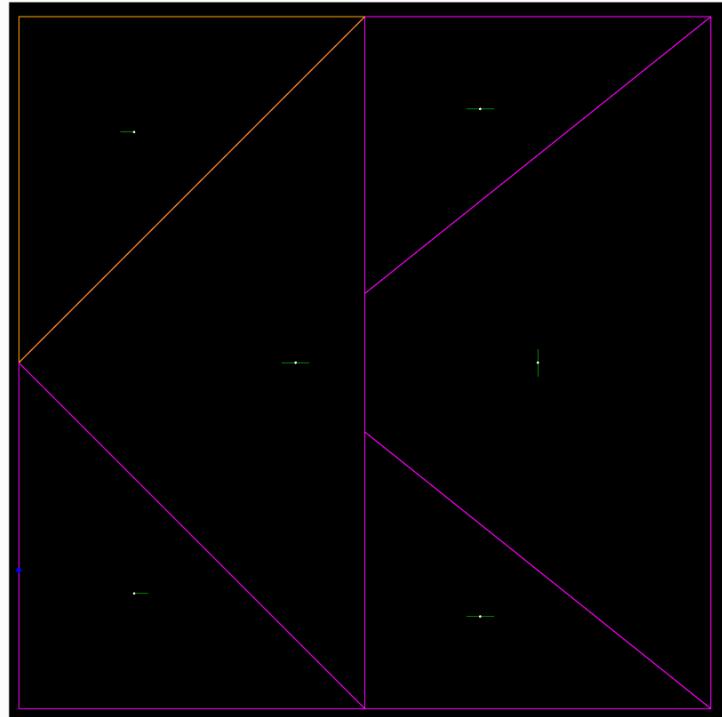
most answers to Q1: decision trees or finite state controllers

most answers to Q2: extract from a neural policy

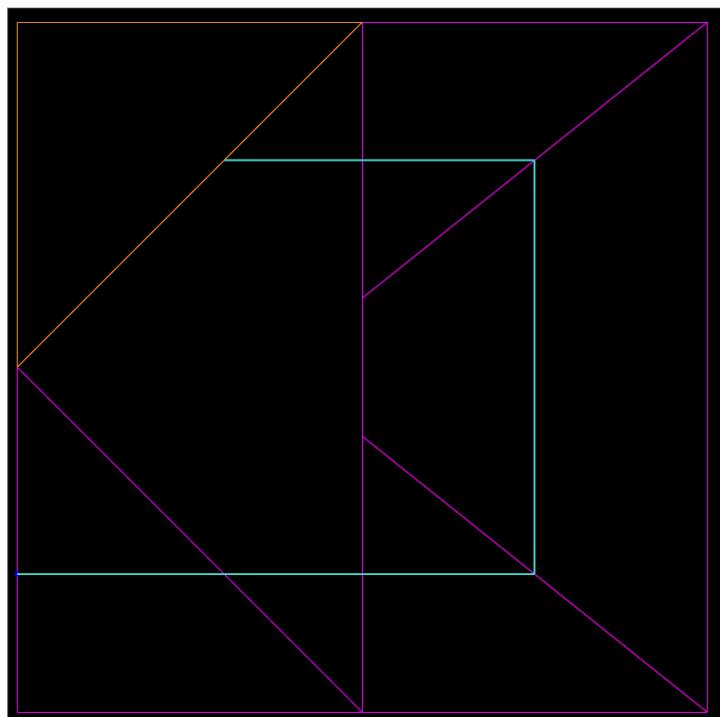
GRIDWORLDS

(Guruprerna Shabati)

problem

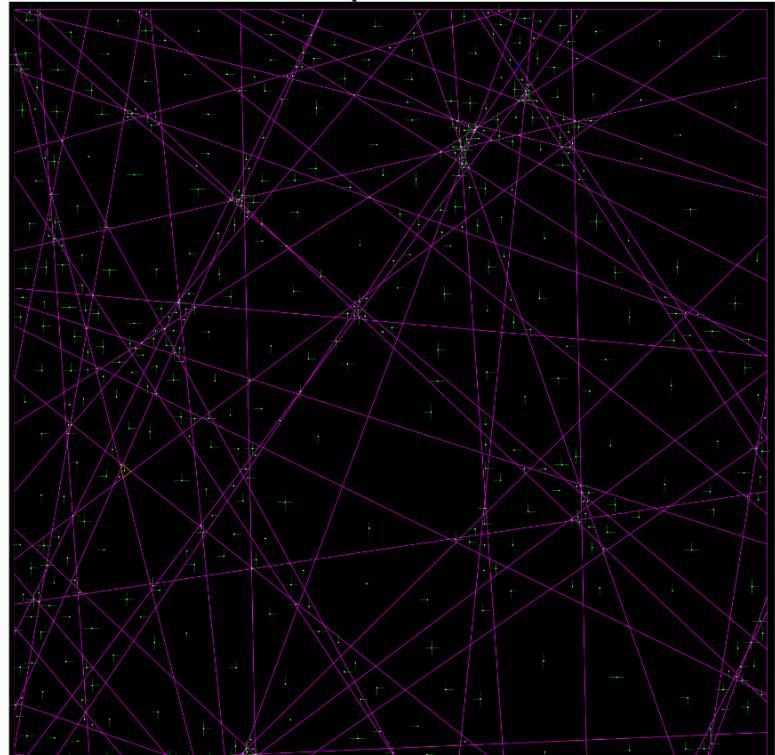


Solution

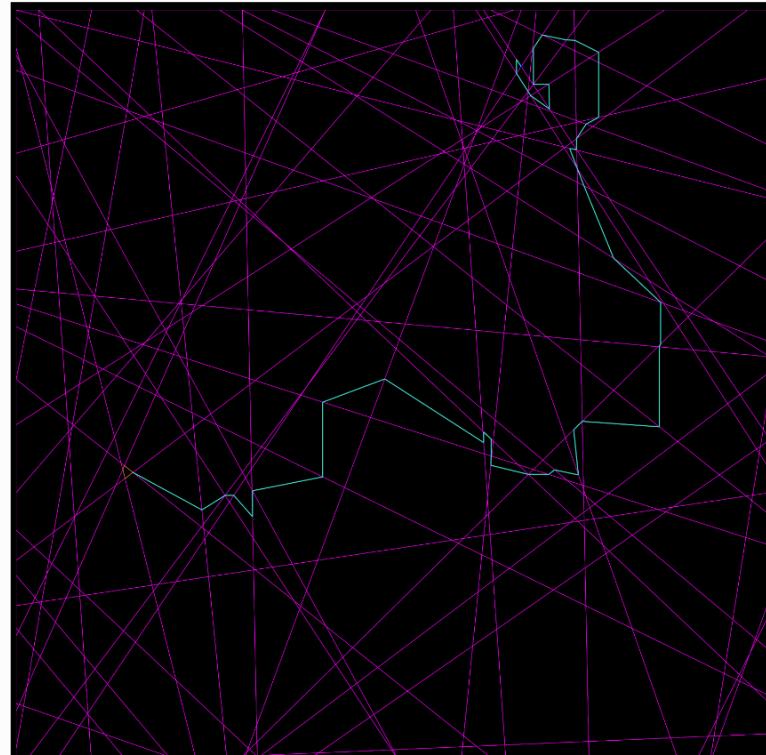


Key feature: The environment is generated as a PRISM program

problem



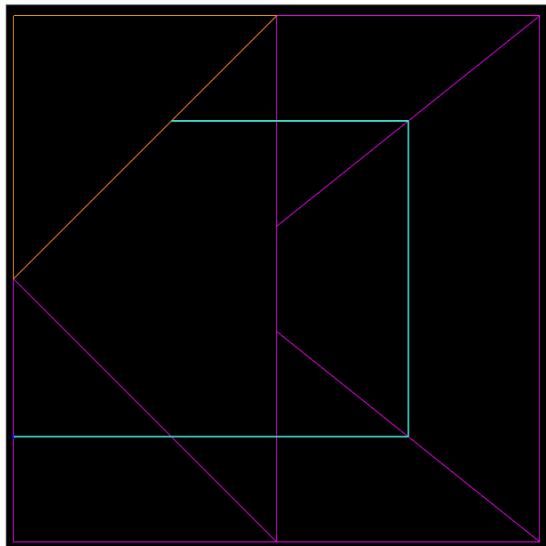
Solution



COMPARING TWO APPROACHES

- (1) Solve as explicit MDP and return optimal policy as explicit function states → actions
- (2) find a policy as program using backward analysis

Example



generated policy

```
Until([(0, 0), (0, 25)]):  
[(0, 25), (0, 0)] ->  
[(5, 20), (25, 0)], Preference: Neutral, Edge: [(0, 25), (25, 0)]  
  
Until([(0, 25), (25, 0)]):  
[(0, 25), (25, 0)] ->  
[(25, 0), (25, 20)], Preference: Neutral, Edge: [(25, 20), (25, 0)]  
  
Until([(25, 0), (25, 20)]):  
[(25, 20), (25, 0)] ->  
[(25, 20), (50, 0)], Preference: Neutral, Edge: [(25, 20), (50, 0)]  
  
Until([(25, 20), (50, 0)]):  
[(25, 20), (50, 0)] ->  
[(25, 30), (50, 50)], Preference: Neutral, Edge: [(25, 30), (50, 50)]  
  
Until([(25, 30), (50, 50)]):  
[(25, 30), (50, 50)] ->  
[(25, 30), (25, 50)], Preference: Neutral, Edge: [(25, 30), (25, 50)]  
  
Until([(25, 30), (25, 50)]):  
[(25, 30), (25, 50)] ->  
[(0, 25), (25, 50)], Preference: Neutral, Edge: [(0, 25), (25, 50)]  
  
Until([(0, 25), (25, 50)]):  
Win!
```

Summary of experimental results:

The program as policies are infinitely more succinct !

and **Verifiable, explainable, readable, executable, ...**

environment as program : **SMALL**

environment as explicit MDP : **LARGE**

policy as program : **SMALL**

policy as explicit S→A : **LARGE**

Motto:

Program you are, and to program you will return

environment

policy

JOIN US

We have open positions in Bordeaux !

PhD and Postdoc

both theory and practice